

```
In [3]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
%matplotlib inline
```

```
In [4]: df=pd.read_csv("Student_Performance.csv")
```

```
In [5]: df.head()
```

Out[5]:

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	Sample Question Papers Practiced	Performance Index
0	7	99	Yes	9	1	91.0
1	4	82	No	4	2	65.0
2	8	51	Yes	7	2	45.0
3	5	52	Yes	5	2	36.0
4	7	75	No	8	5	66.0

```
In [6]: ##check null values
df.isnull().sum()
```

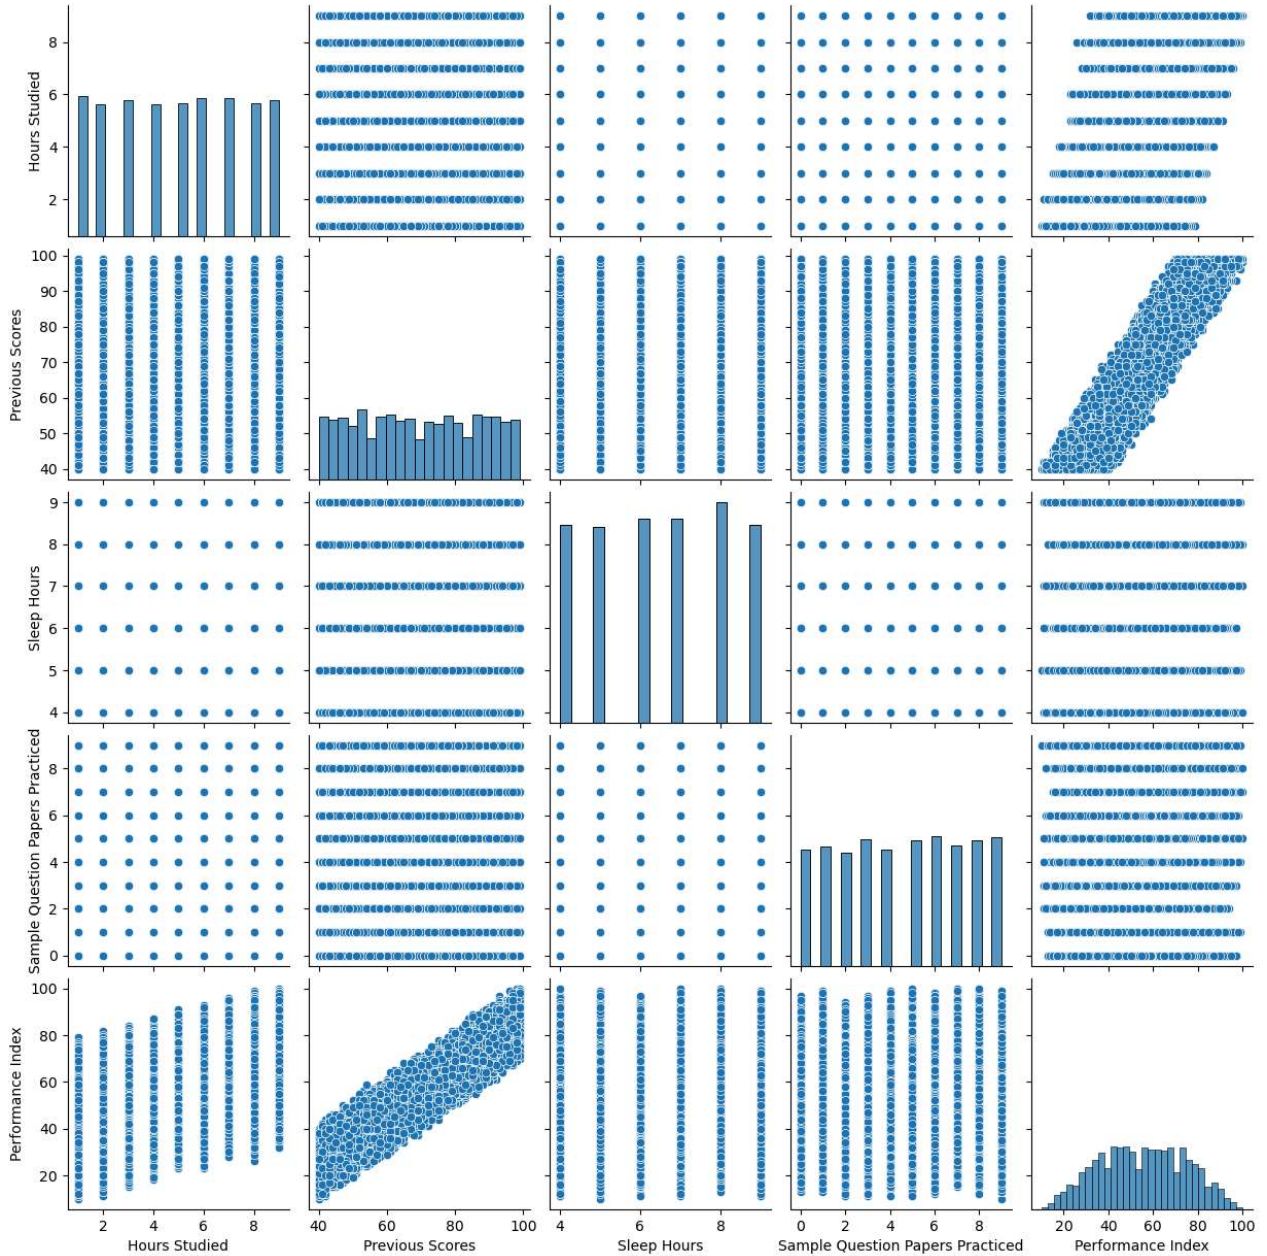
```
Out[6]: Hours Studied      0
Previous Scores           0
Extracurricular Activities 0
Sleep Hours               0
Sample Question Papers Practiced 0
Performance Index         0
dtype: int64
```

```
In [9]: ## Lets do some vidualization
import seaborn as sns
sns.pairplot(df)
```

C:\Users\win 10\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight

```
self._figure.tight_layout(*args, **kwargs)
```

```
Out[9]: <seaborn.axisgrid.PairGrid at 0x195ce978cd0>
```



```
In [94]: df["Extracurricular Activities"]=df["Extracurricular Activities"].map({"Yes":1 , "No":2})
df
```


Out[94]:

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	Sample Question Papers Practiced	Performance Index
0	7	99	1	9	1	91.0
1	4	82	2	4	2	65.0
2	8	51	1	7	2	45.0
3	5	52	1	5	2	36.0
4	7	75	2	8	5	66.0
...
9995	1	49	1	4	2	23.0
9996	7	64	1	8	5	58.0
9997	6	83	1	8	5	74.0
9998	9	97	1	7	0	95.0
9999	7	74	2	8	1	64.0

10000 rows × 6 columns

```
In [95]: ## Visualiza the datapoints more closely
plt.scatter(df['Hours Studied'],df['Performance Index'], color='r')
plt.xlabel("Study Hour")
plt.ylabel("Performance")
```

Out[95]: Text(0, 0.5, 'Performance')



```
In [96]: ##independent and dependent features
X=df.iloc[:, :-1]
y=df.iloc[:, -1]
```

In [97]: `X.head()`

Out[97]:

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	Sample Question Papers Practiced
0	7	99	1	9	1
1	4	82	2	4	2
2	8	51	1	7	2
3	5	52	1	5	2
4	7	75	2	8	5

In [98]: `y`

Out[98]:

```
0      91.0
1      65.0
2      45.0
3      36.0
4      66.0
...
9995   23.0
9996   58.0
9997   74.0
9998   95.0
9999   64.0
Name: Performance Index, Length: 10000, dtype: float64
```


In [112]: `# train test split`
`from sklearn.model_selection import train_test_split`

In [113]: `X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25,random_state=42)`

In [114]: `import seaborn as sns`


In [115]: `sns.regplot(x=df['Sleep Hours'],y=df['Hours Studied'])`

Out[115]: `<Axes: xlabel='Sleep Hours', ylabel='Hours Studied'>`




```
In [116]: sns.regplot(x=df['Sleep Hours'],y=df['Previous Scores'])
```

```
Out[116]: <Axes: xlabel='Sleep Hours', ylabel='Previous Scores'>
```



```
In [117]: sns.regplot(x=df['Performance Index'],y=df['Sample Question Papers Practiced'])
```

```
Out[117]: <Axes: xlabel='Performance Index', ylabel='Sample Question Papers Practiced'>
```



```
In [138]: x=df.drop("Performance Index",axis=1)
           y=df['Performance Index']
           train_X,test_X,train_Y,test_Y = train_test_split(x,y,test_size=0.2,shuffle=True)
```

```
In [139]: from sklearn.linear_model import LinearRegression
```

```
In [140]: regression=LinearRegression()
```

```
In [141]: regression.fit(X_train,y_train)
```

```
Out[141]: ▾ LinearRegression  
LinearRegression()
```

```
In [142]: ## cross validation  
from sklearn.model_selection import cross_val_score
```

```
In [143]: validation_score=cross_val_score(regression,X_train,y_train,scoring='neg_mean_squared_error',cv=3)
```

```
In [144]: np.mean(validation_score)
```

```
Out[144]: -4.2027969495394295
```

```
In [145]: #Prediction  
y_pred=regression.predict(X_test)
```

```
In [146]: y_pred
```

```
Out[146]: array([54.73187888, 22.61211054, 47.90838844, ..., 68.07396952,  
53.68636805, 54.85816372])
```

```
In [147]: ## Performance Metrics  
from sklearn.metrics import mean_absolute_error, mean_squared_error  
mse=mean_squared_error(y_test,y_pred)  
mae=mean_absolute_error(y_test,y_pred)  
rmse=np.sqrt(mse)  
print(mse)  
print(mae)  
print(rmse)
```

```
4.032544215419129  
1.5975792091646137  
2.0081195719924474
```

```
In [148]: from sklearn.metrics import r2_score  
score=r2_score(y_test,y_pred)  
print(score)
```

```
0.9890550757439103
```


```
In [149]: #display adjusted R-squared  
print(1 - (1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1))
```

```
0.9890331332333728
```

Assumptions

```
In [150]: plt.scatter(y_pred,y_test)
```

```
Out[150]: <matplotlib.collections.PathCollection at 0x195dc03da10>
```



```
In [151]: residuals=y_test-y_pred  
print(residuals)
```


```
6252  -3.731879  
4684  -2.612111  
1731  -1.908388  
4742  -3.301042  
4521  -2.035815  
...  
4862   0.575641  
7025  -1.000419  
7647   1.926030  
7161   2.313632  
73    -2.858164  
Name: Performance Index, Length: 2500, dtype: float64
```



```
In [152]: ## Plot this residuals  
sns.displot(residuals, kind='kde')
```


C:\Users\win 10\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)

Out[152]: <seaborn.axisgrid.FacetGrid at 0x195da62fc10>




```
In [153]: ## scatter plot with respect to prediction and residuals  
plt.scatter(y_pred,residuals)
```

```
Out[153]: <matplotlib.collections.PathCollection at 0x195dc0f8390>
```



```
In [154]: ## OLS Linear Regression  
import statsmodels.api as sm  
model=sm.OLS(y_train,X_train).fit()
```

In [135]:

model.summary()

Out[135]:

OLS Regression Results

Dep. Variable:	Performance Index	R-squared (uncentered):	0.992
Model:	OLS	Adj. R-squared (uncentered):	0.992
Method:	Least Squares	F-statistic:	1.890e+05
Date:	Mon, 02 Jun 2025	Prob (F-statistic):	0.00
Time:	22:44:13	Log-Likelihood:	-23011.
No. Observations:	7500	AIC:	4.603e+04
Df Residuals:	7495	BIC:	4.607e+04
Df Model:	5		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Hours Studied	2.3392	0.022	105.210	0.000	2.296	2.383
Previous Scores	0.8540	0.003	304.681	0.000	0.849	0.860
Extracurricular Activities	-4.7304	0.108	-43.658	0.000	-4.943	-4.518
Sleep Hours	-1.0713	0.030	-36.251	0.000	-1.129	-1.013
Sample Question Papers Practiced	-0.1889	0.020	-9.235	0.000	-0.229	-0.149

Omnibus:	18.746	Durbin-Watson:	1.929
Prob(Omnibus):	0.000	Jarque-Bera (JB):	15.050
Skew:	-0.003	Prob(JB):	0.000539
Kurtosis:	2.781	Cond. No.	131.

Notes:

[1] R² is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [137]:

print(regression.coef_)

[2.85492123 1.01637916 -0.58370931 0.47688351 0.19092346]